

Probabilistic Algorithms

Paul Cotofrei

information management institute

PA 2016

Outline

Motivation and supporting results

Multiple Objective Functions

No Free Lunch theorem

Some problems asking optimization

- ▶ Find the best red-yellow-green signal timings in an urban traffic network
- ▶ Determine the optimal schedule for use of laboratory facilities to serve an organization's overall interests
- ▶ Minimize the costs of shipping from production facilities to warehouses
- ▶ Maximize the probability of detecting an incoming warhead (vs. decoy) in a missile defense system
- ▶ Place sensors in manner to maximize useful information
- ▶ Determine the times to administer a sequence of drugs for maximum therapeutic effect

Model the problems as a mathematical model depending on:

1. a set of adjustable parameters/variables
2. an (objective) function defined on the set of parameters
3. a goal: minimize/maximize the function

Two Fundamental Problems of Interest

- ▶ Let \mathbf{D} be the domain of allowable values for a vector \mathbf{d}
- ▶ \mathbf{d} represents a vector of "adjustables" and may be continuous or discrete (or both)


Two fundamental problems of interest

- ▶ **Problem 1.** Find the value(s) of a vector $\mathbf{d} \in \mathbf{D}$ that minimize a scalar-valued *loss function* $L(\mathbf{d})$
- ▶ **Problem 2.** Find the value(s) of $\mathbf{d} \in \mathbf{D}$ that solve the equation $g(\mathbf{d}) = 0$ for some vector-valued function $g(\mathbf{d})$


Frequently (but not necessarily) $g(\mathbf{d}) = \partial L(\mathbf{d}) / \partial \mathbf{d}$

Convert P1 \Rightarrow P2 ? Convert P2 \Rightarrow P1?


Three Common Types of Loss Functions



Continuous



**Discrete/
Continuous**



Discrete

Forms of loss function

- ▶ Mathematical expression
- ▶ Algorithms (simulations)
- ▶ Physical experiments

Classical Calculus-Based Optimization

- ▶ Classical optimization setting of interest

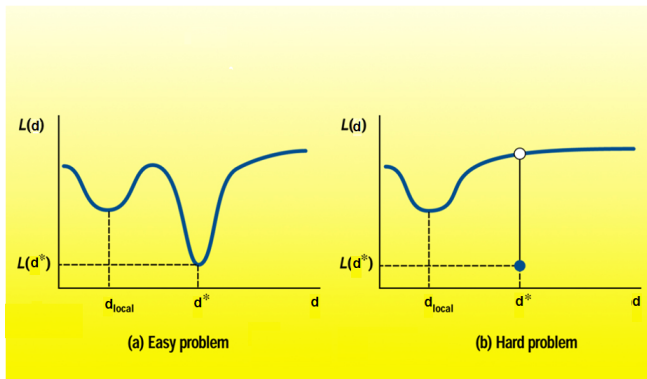
$$\mathbf{D}^* \equiv \min_{\mathbf{d} \in \mathbf{D}} L(\mathbf{d}) = \{\mathbf{d}^* \in \mathbf{D} : L(\mathbf{d}^*) \leq L(\mathbf{d}) \text{ for all } \mathbf{d} \in \mathbf{D}\}$$

where \mathbf{d} is a n -dimensional vector of parameters and $\mathbf{D} \in \mathcal{R}^n$ is the domain representing the constraints on allowable values for \mathbf{d} .

- ▶ \mathbf{D}^* : a single point, a countable collection of points or an uncountable number
 1. $L(\mathbf{d}) = \mathbf{d}^T \mathbf{d}$, $\mathbf{D} = \mathcal{R}^n$; $\mathbf{D}^* = ?$
 2. $L(\mathbf{d}) = \sin(\mathbf{d})$, $\mathbf{D} = [0, 4\pi]$; $\mathbf{D}^* = ?$
 3. $L(\mathbf{d}) = \cos(\mathbf{d})$, $\mathbf{D} = \mathcal{R}$; $\mathbf{D}^* = ?$
 4. $L(\mathbf{d}) = (\mathbf{d}^T \mathbf{d} - 1)^2$, $\mathbf{D} = \mathcal{R}^n$; $\mathbf{D}^* = ?$

Global vs. Local Solutions

- ▶ Any $\mathbf{d}^* \in \mathbf{D}^*$ is a *global* solution
- ▶ A *local* solution \mathbf{d}_{local} satisfies $L(\mathbf{d}_{local}) \leq L(\mathbf{d})$ for any \mathbf{d} in a vicinity of \mathbf{d}_{local}



Global vs. Local Methods

- ▶ General global optimization problem is very difficult
- ▶ Sometimes local optimization is "good enough" given limited resources available
- ▶ Global methods include: genetic algorithms, evolutionary strategies, simulated annealing, etc.
- ▶ Global methods tend to have following characteristics:
 - ▶ Inefficient, especially for high-dimensional \mathbf{d}
 - ▶ Relatively difficult to use (e.g., require very careful selection of algorithm coefficients)
 - ▶ Sometimes questionable theoretical foundation for global convergence
 - ▶ Multiple runs usually required to have confidence in reaching global optimum

Stochastic Optimization

A. Random noise in input information (e.g., measurements with noise for $L(\mathbf{d})$ or $g(\mathbf{d})$):

$$y(\mathbf{d}) \equiv L(\mathbf{d}) + \varepsilon(\mathbf{d})$$

$$Y(\mathbf{d}) \equiv g(\mathbf{d}) + e(\mathbf{d}),$$

where ε and e represent the noise terms.

B. Injected randomness (Monte Carlo) in choice of algorithm iteration magnitude/direction

- ▶ Contrasts with deterministic methods (e.g., steepest descent, Newton-Raphson, etc.)
 - ▶ Assume perfect information about $L(\mathbf{d})$ (and its gradients)
 - ▶ Search magnitude/direction deterministic at each iteration
- ▶ Injected randomness (B) in search magnitude/direction can offer benefits in efficiency and robustness
 - ▶ E.g., Capabilities for global (vs. local) optimization

General Structure of Optimization Process

Concepts and Definitions

- ▶ **Problem Space:** the set \mathbf{D} containing all elements \mathbf{d} which could be the solution of an optimization problem related to a *loss function* L .
 - ▶ For the same optimization problem, different problem spaces can be defined
 - ▶ The problem space can be restricted by *logical constraints* or *practical constraints*
- ▶ **Solution candidate:** an element $\hat{\mathbf{d}}$ of the problem space \mathbf{D} for a certain optimization problem.
- ▶ **Solution space:** the set \mathbf{D}^* of all solutions of an optimization problem.
- ▶ **Search space:** the set \mathcal{G} of all elements $\hat{\mathbf{d}} \in \mathbf{D}$ which can be processed by an optimization algorithm in order to solve a given problem.
- ▶ **Search operations:** the operations used by optimization algorithms in order to explore the search space \mathcal{G} .

Outline

Motivation and supporting results

Multiple Objective Functions

No Free Lunch theorem

Multiple Objective Functions


- ▶ For many real-world design or decision making problems: optimize the set F of m criterions:

$$F = \{L_i : \mathbf{D}_i \mapsto Y_i : 0 < i \leq m, Y_i \subseteq \mathbb{R}\}$$

- ▶ Example:

1. Maximize profit and Minimize costs for advertising, personal, raw materials etc..
2. Maximize product quality and Minimize negative impact on environment.

- ▶ Minimize $L_1 : \mathbb{R}^2 \mapsto \mathbb{R}$ and $L_2 : \mathbb{R}^2 \mapsto \mathbb{R}$



Two functions L_1 and L_2 with different minima x_1 , x_2 , x_3 , and x_4 .

Approaches for optimum definition


- ▶ **Weighted Sums** (or Linear Aggregation)

$$\text{Minimize } L(\mathbf{d}) = \sum_{i=1}^n \omega_i w_i L_i(\mathbf{d})$$

where $\omega_i = \begin{cases} 1 & \text{if } L_i \text{ should be minimized} \\ -1 & \text{if } L_i \text{ should be maximized} \end{cases}$


and w_i is the relative importance of L_i

- ▶ Example: Consider equal relative importance w_1 and w_2 for L_1 and L_2



Drawbacks of linear aggregation

- ▶ How to determine the weights w_i ?
- ▶ Not suitable for functions from different O classes.
 - ▶ Example: $f_1(x) = x^2$ and $f_2(x) = e^{x-2}$, so $f_1(x) = O(x^2) \neq O(e^x) = f_2(x)$
 - ▶ For x around 0, f_2 is negligible compared to f_1 ; for $x > 5$, f_1 is negligible compared to f_2 .



Pareto Optimization

- It is based on the concept of **Pareto domination**: an element \mathbf{d}_1 dominates (is preferred to) an element \mathbf{d}_2 (i.e. $\mathbf{d}_1 \vdash \mathbf{d}_2$) if \mathbf{d}_1 is better than \mathbf{d}_2 in at least one objective function and not worse with respect to all other objectives.

$$\mathbf{d}_1 \vdash \mathbf{d}_2 \Leftrightarrow \begin{cases} \forall i \in 1..n, \omega_i L_i(\mathbf{d}_1) \leq \omega_i L_i(\mathbf{d}_2), \text{ and} \\ \exists j \in 1..n : \omega_j L_j(\mathbf{d}_1) < \omega_j L_j(\mathbf{d}_2) \end{cases}$$


where

$$\omega_i = \begin{cases} 1 & \text{if } L_i \text{ should be minimized} \\ -1 & \text{if } L_i \text{ should be maximized} \end{cases}$$

- **Pareto optimal**: An element $\mathbf{d} \in \mathbf{D}$ is Pareto optimal if it is not dominated by any other element in the problem space \mathbf{D} .
- The optimal set $\mathbf{D}^* = \{\mathbf{d}^* \in \mathbf{D} \mid \nexists \mathbf{d} \in \mathbf{D} : \mathbf{d} \vdash \mathbf{d}^*\}$

Pareto Optimal


Example: maximize $f_1(x)$ and $f_2(x)$ on $\mathbf{D} = [0, \infty)$



- ▶ $f_1(x_2) > f_1(x)$ and $f_2(x_2) > f_2(x)$ for all $x \in [x_1, x_2)$, so $x_2 \vdash x \forall x \in [x_1, x_2)$.
- ▶ The points $x \in [x_2, x_3]$ are not dominated by any other points, so are Pareto optimal.
- ▶ $f_1(x_5) > f_1(x)$ and $f_2(x_5) > f_2(x)$ for all $x \in [x_3, x_4)$, so $x_5 \vdash x \forall x \in [x_3, x_4)$.
- ▶ The points $x \in [x_5, x_6]$ are not dominated by any other points, so are Pareto optimal.
- ▶ The set \mathbf{D}^* of Pareto optimal points is $[x_2, x_3] \cup [x_5, x_6]$

External Decision Maker

- ▶ A weakness of Pareto optimization: one may have two elements, \mathbf{d}_1 and \mathbf{d}_2 such that neither $\mathbf{d}_1 \vdash \mathbf{d}_2$ nor $\mathbf{d}_2 \vdash \mathbf{d}_1$.
- ▶ For many optimization problems we need a total order: the solution \mathbf{d}_1 is better, equal or worse than solution \mathbf{d}_2
- ▶ A total order using Pareto optimization : *Pareto ranking*
 - ▶ In the first step, the elements not dominated receive rank 0
 - ▶ In the following steps, one removes the elements with rank i from \mathbf{D} , and the new non-dominated elements receive rank $i + 1$
- ▶ The **External Decision Maker**: it uses a-priori knowledge and user preference to provide a cost function $u : \mathbf{Y} \mapsto \mathbb{R}$ which maps the space of objective values to the space of real numbers (total order).



Outline

Motivation and supporting results

Multiple Objective Functions

No Free Lunch theorem

No Free Lunch Theorems

- ▶ Wolpert and Macready (1997) establish several "No Free Lunch" (NFL) Theorems for optimization
- ▶ NFL Theorems apply to settings where parameter set \mathbf{D} and set of loss function values are finite, discrete sets
 - ▶ Relevant for continuous \mathbf{d} problem when considering digital computer implementation
 - ▶ Results are valid for deterministic and stochastic settings
- ▶ Number of optimization problems - mappings from \mathbf{D} to set of loss values - is finite
- ▶ NFL Theorems state, in essence, that no one **optimisation algorithm** is "best" for all problems

No Free Lunch Theorems - Basic Formulation

- ▶ Suppose that

$N_{\mathbf{d}}$ = number of values of \mathbf{d}

N_L = number of values of loss function

- ▶ Then

$(N_L)^{N_{\mathbf{d}}} = \text{number of loss functions}$

- ▶ There is a finite (but possibly huge) number of loss functions
- ▶ Performance measure: the lowest value $L(\hat{\theta})$ obtained after k distinct evaluations of loss function L (i.e. $L(\hat{\mathbf{d}}_1), \dots, L(\hat{\mathbf{d}}_k)$)
- ▶ Basic form of NFL considers average performance over all loss functions

Illustration of No Free Lunch Theorems

- ▶ Three values of \mathbf{d} , two outcomes for noise free loss L
 - ▶ Eight possible mappings, hence eight optimization problems
- ▶ Mean loss across all problems is same regardless of \mathbf{d} ; entries 1 or 2 in table below represent two possible L outcomes

$\mathbf{d} \backslash \text{Map}$	1	2	3	4	5	6	7	8
\mathbf{d}_1	1	1	1	2	2	2	1	2
\mathbf{d}_2	1	1	2	1	1	2	2	2
\mathbf{d}_3	1	2	2	1	2	1	1	2

Overall Consequences of NFL Theorems

- ▶ NFL Theorems state, in essence, that

Averaging (uniformly) over all possible problems (loss functions L), all algorithms perform equally well

- ▶ In particular, if algorithm 1 performs better than algorithm 2 over some set of problems, then algorithm 2 performs better than algorithm 1 on another set of problems

Overall relative efficiency of two algorithms cannot be inferred from a few sample problems

- ▶ NFL theorems say nothing about **specific** algorithms on **specific** problems